

Data Project - Stock Market Analysis



```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid')
plt.style.use("fivethirtyeight")
%matplotlib inline
```

```
In [2]: df = pd.read_csv(r"C:\Users\HP\stocks.csv")
df
```

```
Out[2]:
```

	date	open	high	low	close	volume	Name
0	2013-02-08	15.07	15.12	14.63	14.75	8407500	AAL
1	2013-02-11	14.89	15.01	14.26	14.46	8882000	AAL
2	2013-02-12	14.45	14.51	14.10	14.27	8126000	AAL
3	2013-02-13	14.30	14.94	14.25	14.66	10259500	AAL
4	2013-02-14	14.94	14.96	13.16	13.99	31879900	AAL
...
619035	2018-02-01	76.84	78.27	76.69	77.82	2982259	ZTS
619036	2018-02-02	77.53	78.12	76.73	76.78	2595187	ZTS
619037	2018-02-05	76.64	76.92	73.18	73.83	2962031	ZTS
619038	2018-02-06	72.74	74.56	72.13	73.27	4924323	ZTS
619039	2018-02-07	72.70	75.00	72.69	73.86	4534912	ZTS

619040 rows × 7 columns


```
In [3]: print(df.head())
```

	date	open	high	low	close	volume	Name
0	2013-02-08	15.07	15.12	14.63	14.75	8407500	AAL
1	2013-02-11	14.89	15.01	14.26	14.46	8882000	AAL
2	2013-02-12	14.45	14.51	14.10	14.27	8126000	AAL
3	2013-02-13	14.30	14.94	14.25	14.66	10259500	AAL
4	2013-02-14	14.94	14.96	13.16	13.99	31879900	AAL

In [4]: `print(df.columns)`

Index(['date', 'open', 'high', 'low', 'close', 'volume', 'Name'], dtype='object')

In [5]: `plt.figure(figsize=(16,6))`
`plt.title('Close Price History')`
`plt.plot(df['close'])`
`plt.xlabel('Date', fontsize=18)`
`plt.ylabel('Close Price USD ($)', fontsize=18)`
`plt.show()`



In [6]: `# Create a new dataframe with only the 'Close column`
`data = df.filter(['close'])`
`# Convert the dataframe to a numpy array`
`dataset = data.values`
`# Get the number of rows to train the model on`
`training_data_len = int(np.ceil(len(dataset) * .95))`

`training_data_len`

Out[6]: 588088

In [7]: `# Scale the data`
`from sklearn.preprocessing import MinMaxScaler`

`scaler = MinMaxScaler(feature_range=(0,1))`
`scaled_data = scaler.fit_transform(dataset)`

`scaled_data`

Out[7]: `array([[0.00642763],`
 `[0.00628599],`
 `[0.00619319],`
 `...,`
 `[0.0352836],`
 `[0.03501009],`
 `[0.03529825]])`

In [8]: `# Create the training data set`
`# Create the scaled training data set`
`train_data = scaled_data[0:int(training_data_len), :]`
`# Split the data into x_train and y_train data sets`

```
x_train = []
y_train = []
```

```
for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
    y_train.append(train_data[i, 0])
    if i <= 61:
        print(x_train)
        print(y_train)
        print()

# Convert the x_train and y_train to numpy arrays
x_train, y_train = np.array(x_train), np.array(y_train)

# Reshape the data
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))
# x_train.shape
```

```
[array([0.00642763, 0.00628599, 0.00619319, 0.00638367, 0.00605643,
        0.00630553, 0.00618831, 0.00573407, 0.00575361, 0.0058513 ,
        0.00558266, 0.00569988, 0.00577315, 0.00578292, 0.00587083,
        0.00601247, 0.00608574, 0.00633972, 0.00646182, 0.00651066,
        0.00661323, 0.00679395, 0.0069942 , 0.00716027, 0.00702839,
        0.0071798 , 0.00741913, 0.00763892, 0.00752658, 0.0074582 ,
        0.00733121, 0.00728726, 0.00735563, 0.00751193, 0.0073654 ,
        0.00691117, 0.0068086 , 0.00688675, 0.0069014 , 0.00696001,
        0.00689163, 0.00693071, 0.00713096, 0.00710654, 0.00683791,
        0.00721888, 0.00729214, 0.007087 , 0.00704793, 0.00680372,
        0.00718469, 0.00725795, 0.00714561, 0.00732633, 0.00743378,
        0.00747774, 0.00733121, 0.00749728, 0.00753635, 0.00752658])]
[0.007516813925886851]
```

```
[array([0.00642763, 0.00628599, 0.00619319, 0.00638367, 0.00605643,
        0.00630553, 0.00618831, 0.00573407, 0.00575361, 0.0058513 ,
        0.00558266, 0.00569988, 0.00577315, 0.00578292, 0.00587083,
        0.00601247, 0.00608574, 0.00633972, 0.00646182, 0.00651066,
        0.00661323, 0.00679395, 0.0069942 , 0.00716027, 0.00702839,
        0.0071798 , 0.00741913, 0.00763892, 0.00752658, 0.0074582 ,
        0.00733121, 0.00728726, 0.00735563, 0.00751193, 0.0073654 ,
        0.00691117, 0.0068086 , 0.00688675, 0.0069014 , 0.00696001,
        0.00689163, 0.00693071, 0.00713096, 0.00710654, 0.00683791,
        0.00721888, 0.00729214, 0.007087 , 0.00704793, 0.00680372,
        0.00718469, 0.00725795, 0.00714561, 0.00732633, 0.00743378,
        0.00747774, 0.00733121, 0.00749728, 0.00753635, 0.00752658]), array([0.006285
99, 0.00619319, 0.00638367, 0.00605643, 0.00630553,
        0.00618831, 0.00573407, 0.00575361, 0.0058513 , 0.00558266,
        0.00569988, 0.00577315, 0.00578292, 0.00587083, 0.00601247,
        0.00608574, 0.00633972, 0.00646182, 0.00651066, 0.00661323,
        0.00679395, 0.0069942 , 0.00716027, 0.00702839, 0.0071798 ,
        0.00741913, 0.00763892, 0.00752658, 0.0074582 , 0.00733121,
        0.00728726, 0.00735563, 0.00751193, 0.0073654 , 0.00691117,
        0.0068086 , 0.00688675, 0.0069014 , 0.00696001, 0.00689163,
        0.00693071, 0.00713096, 0.00710654, 0.00683791, 0.00721888,
        0.00729214, 0.007087 , 0.00704793, 0.00680372, 0.00718469,
        0.00725795, 0.00714561, 0.00732633, 0.00743378, 0.00747774,
        0.00733121, 0.00749728, 0.00753635, 0.00752658, 0.00751681])]
[0.007516813925886851, 0.00769264583058596]
```

```
In [9]: from keras.models import Sequential
        from keras.layers import Dense, LSTM

# Build the LSTM model
model = Sequential()
```

```

model.add(LSTM(128, return_sequences=True, input_shape= (x_train.shape[1], 1)))
model.add(LSTM(64, return_sequences=False))
model.add(Dense(25))
model.add(Dense(1))

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Train the model
model.fit(x_train, y_train, batch_size=100, epochs=5)

```

```

Epoch 1/5
5881/5881 [=====] - 1645s 278ms/step - loss: 2.1789e-05
Epoch 2/5
5881/5881 [=====] - 1302s 221ms/step - loss: 9.4617e-06
Epoch 3/5
5881/5881 [=====] - 1422s 242ms/step - loss: 8.9023e-06
Epoch 4/5
5881/5881 [=====] - 1335s 227ms/step - loss: 8.9060e-06
Epoch 5/5
5881/5881 [=====] - 2126s 362ms/step - loss: 8.2106e-06
Out[9]: <keras.src.callbacks.History at 0x232a08ba990>

```

```

In [10]: # Create the testing data set
# Create a new array containing scaled values from index 1543 to 2002
test_data = scaled_data[training_data_len - 60: , :]
# Create the data sets x_test and y_test
x_test = []
y_test = dataset[training_data_len:, :]
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:i, 0])

# Convert the data to a numpy array
x_test = np.array(x_test)

# Reshape the data
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1 ))

# Get the models predicted price values
predictions = model.predict(x_test)
predictions = scaler.inverse_transform(predictions)

# Get the root mean squared error (RMSE)
rmse = np.sqrt(np.mean(((predictions - y_test) ** 2)))
rmse

```

```

968/968 [=====] - 63s 50ms/step
2.103550450280048

```

Out[10]:


```

In [11]: # Plot the data
train = data[:training_data_len]
valid = data[training_data_len:]
valid['Predictions'] = predictions
# Visualize the data
plt.figure(figsize=(16,6))
plt.title('Model')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train['close'])
plt.plot(valid[['close', 'Predictions']])
plt.legend(['Train', 'Val', 'Predictions'], loc='lower right')
plt.show()

```

C:\Users\HP\AppData\Local\Temp\ipykernel_5396\2269485014.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
valid['Predictions'] = predictions



In [12]: `# Show the valid and predicted prices`
`valid`

Out[12]:

	close	Predictions
588088	68.47	67.932472
588089	68.61	68.589119
588090	67.77	68.735733
588091	66.75	67.926224
588092	67.38	66.906975
...
619035	77.82	76.845627
619036	76.78	77.870880
619037	73.83	76.809448
619038	73.27	73.946838
619039	73.86	73.352760

30952 rows × 2 columns

Summary

In this notebook, you discovered and explored stock data.

Specifically, you learned:

How to load stock market data from the YAHOO Finance website using yfinance. How to explore and visualize time-series data using Pandas, Matplotlib, and Seaborn. How to measure the correlation between stocks. How to measure the risk of investing in a particular

stock. Do you have any questions? Ask your questions in the comments below and I will do my best to answer.

References: <https://www.investopedia.com/terms/c/correlation.asp> Jose Portilla Udemy
Course: Learning Python for Data Analysis and Visualization

